

Forensic Analysis Report

Neuro-Symbolic Crisis Generator • Comprehensive Evaluation

Generated: 2025-12-18 17:57:07

Test Configuration

Model: GPT-4o

Temperature: 0.7

Mode: Thesis (Full Validation) vs. Legacy (Skip Validation)

Max Iterations: 20 per Scenario

Execution: Parallel (2 Scenarios simultaneously)

Executive Summary

Note: Limited batch comparison data available. For detailed analysis, see Forensic Analysis section below.

2

TOTAL SCENARIOS

0

HALLUCINATIONS PREVENTED

0.00

AVG LEGACY HALLUCINATIONS

0.00

AVG THESIS HALLUCINATIONS

0.0%

AVERAGE REDUCTION

A/B Testing Comparison

A/B Testing Comparison



Statistical Significance Analysis

Hypothesis Testing

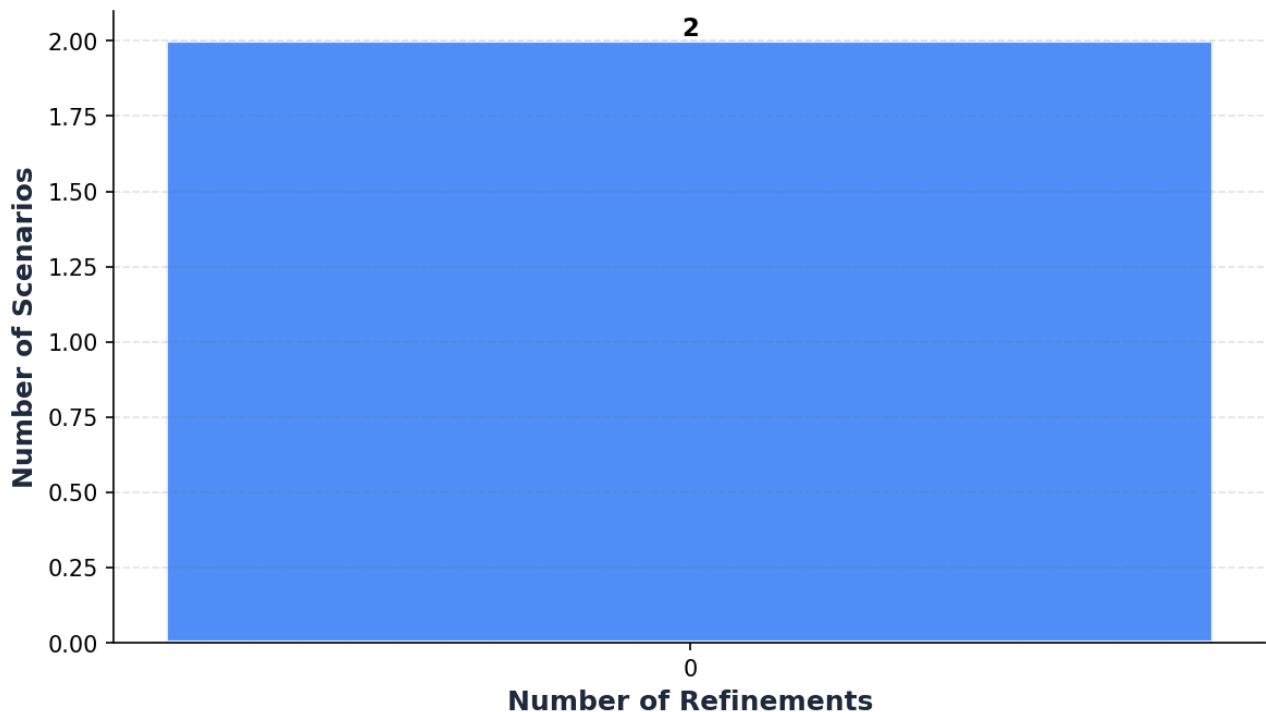
H0: $\mu_{\text{legacy}} = \mu_{\text{thesis}}$ (No difference)

H1: $\mu_{\text{legacy}} > \mu_{\text{thesis}}$ (Thesis reduces hallucinations)

Metric	Value	Interpretation
Mean Difference	0.000	Legacy - Thesis hallucinations
T-Statistic	nan	Test statistic
P-Value	NAN	Significance level
Cohen's d	0.000	Effect size: Negligible
95% CI	[0.000, 0.000]	Confidence interval
Sample Size	2	Number of scenarios

Refinement Efficiency Analysis

Refinement Distribution



0.00

AVG REFINES PER SCENARIO

0

MAX REFINES NEEDED

100.0%

FIRST ATTEMPT SUCCESS RATE

Cost-Benefit Analysis

54

Avg Legacy API Calls

54

Avg Thesis API Calls

0.0%

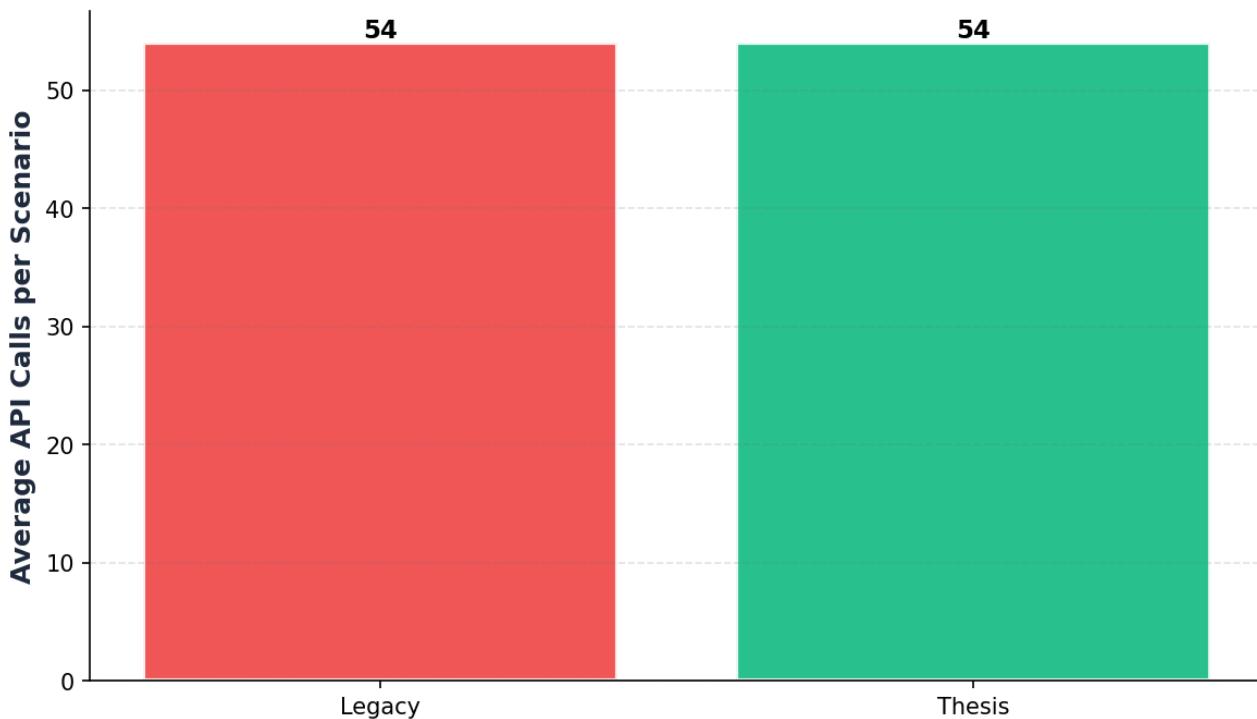
Overhead

0

Total Prevented

API Call Efficiency

API Call Efficiency Comparison



Detailed Results by Scenario

SCENARIO ID	LEGACY HALLUCINATIONS	THESIS HALLUCINATIONS	HALLUCINATIONS PREVENTED	THESIS REFINES
SCEN-000	0	0	0	0
SCEN-001	0	0	0	0